

# ESBMC-GPU

## A Context-Bounded Model Checking Tool to Verify CUDA Programs

Felipe R. Monteiro<sup>1</sup>, Erickson H. da S. Alves<sup>1,2</sup>, Isabela S. Silva<sup>1</sup>,  
Hussama I. Ismail<sup>1</sup>, Lucas C. Cordeiro<sup>1,3</sup>, and Eddie B. de Lima Filho<sup>1,4</sup>

<sup>1</sup>*Faculty of Technology, Federal University of Amazonas, Brazil*

<sup>2</sup>*Samsung Research Institute, Brazil*

<sup>3</sup>*Department of Computer Science, University of Oxford, United Kingdom*

<sup>4</sup>*TPV Technology Limited, Brazil*

---

### Abstract

The Compute Unified Device Architecture (CUDA) is a programming model used for exploring the advantages of graphics processing unit (GPU) devices, through parallelization and specialized functions and features. Nonetheless, as in other development platforms, errors may occur, due to traditional software creation processes, which may even compromise the execution of an entire system. In order to address such a problem, ESBMC-GPU was developed, as an extension to the Efficient SMT-Based Context-Bounded Model Checker (ESBMC). In summary, ESBMC processes input code through ESBMC-GPU and an abstract representation of the standard CUDA libraries, with the goal of checking a set of desired properties. Experimental results showed that ESBMC-GPU was able to correctly verify 85% of the chosen benchmarks and it also overcame other existing GPU verifiers regarding the verification of data-race conditions, array out-of-bounds violations, assertive statements, pointer safety, and the use of specific CUDA features.

*Keywords:* GPU verification, formal verification, model checking, CUDA

---

### 1. Introduction

The Compute Unified Device Architecture (CUDA) is a development framework that makes use of the architecture and processing power of graphics processing units (GPUs) [1]. Indeed, CUDA is also an application programming interface (API), through which a GPU's parallelization scheme and tools can be accessed, with the goal of executing kernels [1]. Nonetheless, source code is still written by human programmers, which may result in

8 arithmetic overflow, division by zero, and other violation types. In addition,  
9 given that CUDA allows parallelization, problems related to the latter can  
10 also occur, due to thread scheduling [2].

11 In order to address the mentioned issues, an extension to the Efficient  
12 SMT-Based Context-Bounded Model Checker (ESBMC) [3] was developed,  
13 named as ESBMC-GPU [4, 5, 6], with the goal of verifying CUDA-based pro-  
14 grams (available online at <http://esbmc.org/gpu/>). ESBMC-GPU consists  
15 of an extension for parsing CUDA source code (*i.e.*, a front-end to ESBMC)  
16 and a CUDA operational model (COM), which is an abstract representation  
17 of the standard CUDA libraries (*i.e.*, the native API) that conservatively  
18 approximates their semantics.

19 A distinct feature of ESBMC-GPU, when compared with other approaches  
20 [2, 7, 8, 9], is the use of Bounded Model Checking (BMC) [10] allied to  
21 Satisfiability Modulo Theories (SMT) [11], with explicit state-space explo-  
22 ration [12, 3]. In summary, concurrency problems are tackled, up to a given  
23 loop/recursion unwinding and context bound, while each interleaving itself  
24 is symbolically handled; however, even with BMC, state-space exploration  
25 may become a very time-consuming task, which is alleviated through state  
26 hashing and Monotonic Partial Order Reduction (MPOR) [13]. As a conse-  
27 quence, redundant interleavings are eliminated, without ignoring a program’s  
28 behavior.

29 Finally, existing GPU verifiers often ignore some aspects related to mem-  
30 ory leak, data transfer, and overflow, which are normally present in CUDA  
31 programs. The proposed approach, in turn, explicitly addresses them, through  
32 an accurate checking procedure, which even considers data exchange between  
33 main program and kernel. Obviously, it results in higher verification times,  
34 but more errors can then be identified and later corrected, in another devel-  
35 opment cycle.

36  
37 **Existing GPU Verifiers.** In addition to ESBMC-GPU, there are other  
38 tools able to verify CUDA programs and each one of them uses its own ap-  
39 proach and targets specific property violations. For instance, GPUVerify [2] is  
40 based on synchronous, delayed visibility semantics, which focuses on detect-  
41 ing data race and barrier divergence, while reducing kernel verification proce-  
42 dures for the analysis of sequential programs. GPU+KLEE (GKLEE) [8], in  
43 turn, is a concrete plus symbolic execution tool, which considers both *kernels*  
44 and *main* functions, while checking deadlocks, memory coalescing, data race,  
45 warp divergence, and compilation level issues. In addition, Concurrency In-  
46 termediate Verification Language (CIVL) [9], a framework for static analysis  
47 and concurrent program verification, uses abstract syntax tree and partial  
48 order reduction to detect user-specified assertions, deadlocks, memory leaks,  
49 invalid pointer dereference, array out-of-bounds, and division by zero.

50 In fact, ESBMC-GPU differs from the aforementioned approaches due to  
 51 its combination of techniques to prune the state-space exploration (*i.e.*, two-  
 52 thread analysis, state hashing, and MPOR) with COM, which demonstrated  
 53 effectiveness in the verification of data-race conditions, array out-of-bounds  
 54 violations, assertive statements, pointer safety, and the use of specific CUDA  
 55 features (cf. Section 5).

## 56 2. Architecture and Implementation

57 ESBMC-GPU is built on top of ESBMC, which is an open-source context-  
 58 bounded model checker based on SMT solvers for ANSI-C/C++ programs [12,  
 59 3, 14], and adds four essential models, as described below.

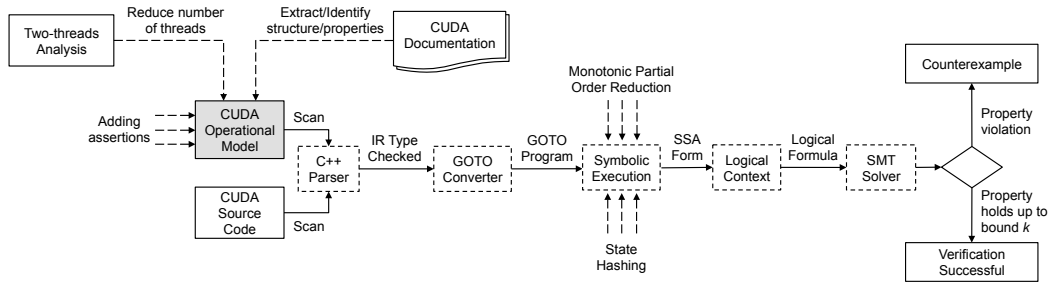


Figure 1: Overview of ESBMC-GPU’s architecture.

60 **CUDA Operational Model.** An operational model for CUDA libraries  
 61 that provides support to CUDA functionalities, in conjunction with ESBMC,  
 62 is shown in Fig. 1. Such an approach, which was previously attempted in the  
 63 verification of C++ programs [14, 15, 16, 17], consists of an abstract repre-  
 64 sentation that reliably approximates the CUDA library’s semantics; however,  
 65 COM incorporates pre- and post-conditions into verification processes, which  
 66 enables ESBMC-GPU to verify specific properties (cf. Section 3). Indeed,  
 67 COM allows the necessary control for performing code analysis, where both  
 68 CUDA operation and knowledge for model checking its properties are avail-  
 69 able. Importantly, COM encloses a representation for the Runtime, Math,  
 70 and cuRAND APIs, which are important CUDA libraries widely used in real  
 71 applications [1]. In particular, with respect to the number of methods/func-  
 72 tions, COM covers 53%, 31%, and 17% of Runtime, Math, and cuRAND  
 73 APIs, respectively.

74 ESBMC was designed to handle multi-threaded software, through the use  
 75 of an API called POSIX – ISO/IEC 9945 [18]. In order to support the veri-  
 76 fication of CUDA kernels, COM applies code transformations to kernel calls  
 77 using ESBMC’s intrinsic functions [6]. In particular, thread/block configu-  
 78 rations in a CUDA kernel call are used as parameters in such intrinsic func-  
 79 tions, which are responsible for checking for preconditions, configuring block

and threads dimension, and translating GPU threads to POSIX ones. Thus, COM is able to support thread interleaving (to create execution paths) and dynamic creation of threads, in order to check data race and specific C/C++ programming language failures (*e.g.*, array out-of-bounds and pointer safety). ESBMC models the Pthreads API [12] to support thread synchronization, *i.e.*, mutex locking operations and conditional waiting, and dynamic creation of threads, which makes that representation very similar to the official CUDA scheduler [6], in such a way that our multi-threading model approximates to that of GPU kernels. Therefore, COM is able to use such aspects present in ESBMC, in order to handle variables in different types of memory and also in inter-warp communication.

The ESBMC’s memory model uses static pointer analysis, padding in structures, with the goal of making all fields align to word boundaries, memory access alignment rules enforcement, and byte array allocation, when the type of memory allocation is unclear [19]. Since ESBMC-GPU is built on top of ESBMC, its memory model for the different types of memory is complete.

**Two-threads Analysis.** Similarly to GPUVerify [2] and PUG [7], ESBMC-GPU also reduces the number of threads (to only two elements), during the verification of CUDA programs, by considering a NVIDIA Fermi GPU architecture [1], in order to improve verification time and avoid the state-space explosion problem. In CUDA programs, whilst threads execute the same parametrized kernel, only two of them are necessary for conflict check. Thus, such an analysis ensures that errors (*e.g.*, data races) detected between two threads, in a given subgroup and due to unsynchronized accesses to shared variables, are enough to justify a property violation [6].

**State Hashing.** ESBMC-GPU applies state hashing to further eliminate redundant interleavings and also reduce the state space, based on SHA256 hashes [20]. In particular, its symbolic state hashing approach computes a summary for a particular state that has already been explored and then indexes the resulting set, in order to reduce the generation of redundant states. Given any state computed during the symbolic execution of a specific CUDA kernel, ESBMC-GPU simply summarizes it and efficiently determines whether it has been explored before or not, along a different computation path. When this behavior is confirmed, which happens during the ESBMC-GPU’s symbolic-execution procedure, then the current computation path does not need to be further explored in the associated reachability tree (RT). This way, if ESBMC-GPU reaches such a state, *i.e.*, where a context switch can be taken (*e.g.*, before a global variable or synchronization primitive) and all shared/local variables and program counters are similar to another explored node, then ESBMC-GPU just considers that an identical node to be

122 further explored, since reachability subtrees associated to them are also sim-  
123 ilar [6, 21].

124

125 **Monotonic Partial Order Reduction.** MPOR is used to reduce the  
126 number of thread interleavings, by classifying transitions inside a program  
127 as dependent or independent. As a consequence, it is possible to determine  
128 whether interleaving pairs always lead to the same state and then remove  
129 duplicates in a RT, without ignoring any program’s behavior [21].

### 130 3. Functionalities

131 COM models CUDA libraries and provides a multi-threading model much  
132 similar to the CUDA scheduler (cf. Section 2). Moreover, it is able to simu-  
133 late CUDA’s program structure and memory, being susceptible of handling  
134 CUDA programs. Thus, through the integration of COM into ESBMC (*i.e.*,  
135 ESBMC-GPU), one is able to analyze CUDA programs and verify the fol-  
136 lowing properties:

137 **Data race.** ESBMC-GPU checks data race conditions, in order to  
138 detect if multiple threads perform unsynchronized access to the same  
139 memory locations;

140 **Pointer safety.** ESBMC-GPU also ensures that (*i*) a pointer offset  
141 does not exceed object bounds and (*ii*) a pointer is neither NULL nor  
142 invalid;

143 **Array bounds.** ESBMC-GPU performs array-bound checking, in or-  
144 der to ensure that any variable, used as an array index, is within known  
145 bounds;

146 **Arithmetic under- and overflow.** ESBMC-GPU checks whether a  
147 sum or product exceeds the memory limits that a variable can han-  
148 dle, which can cause an error capable of spreading through the entire  
149 execution path;

150 **Division by zero.** ESBMC-GPU analyzes whether denominators, in  
151 arithmetic expressions, lead to divisions by zero;

152 **User-specified assertions.** ESBMC-GPU considers all assertions  
153 specified by users, which is essential to a thorough verification process,  
154 as some specific possible violations must be explicitly pointed out.

155 In order to check the aforementioned properties, ESBMC-GPU explicitly  
156 explores the possible interleavings (up to the given context bound) and calls

157 the single-threaded BMC procedure on each one, whenever it reaches a RT  
158 leaf node. Then, the mentioned procedure will stop if it finds a bug or when  
159 all possible RT interleavings have been systematically explored [6]. Further-  
160 more, ESBMC-GPU has the following additional command-line options:

161     `--no-assertions:` to ignore assertions;  
162     `--no-bounds-check:` to skip array bounds check;  
163     `--no-div-by-zero-check:` to skip division by zero check;  
164     `--no-pointer-check:` to skip pointer check;  
165     `--memory-leak-check:` to enable memory leak check;  
166     `--overflow-check:` to enable arithmetic under- and overflow check;  
167     `--deadlock-check:` to enable global/local deadlock check with mutex;  
168     `--data-races-check:` to enable data races check;  
169     `--lock-order-check:` to enable for lock acquisition ordering check;  
170     `--atomicity-check:` to enable atomicity check at visible assignments;  
171     `--force-malloc-success:` to consider that there is always enough  
172     memory available in the device;

173 Thus, ESBMC-GPU is able to check CUDA programs for: deadlock, asser-  
174 tion, lock acquisition error, division by zero, pointer safety, arithmetic over-  
175 flow, and/or out-of-bounds array violation. The precision and performance  
176 of ESBMC-GPU will be further discussed in Section 5.

#### 177 4. Illustrative Example

178 In this part, ESBMC-GPU usage is demonstrated, by using the CUDA  
179 program shown in Fig. 2. First of all, users must replace the default ker-  
180 nel call (line 16) by an intrinsic function of ESBMC-GPU (line 17). Then,  
181 the resulting CUDA program can be passed to the command-line version of  
182 ESBMC-GPU, as follows:

183     `esbmc-gpu <file>.cu --unwind <k> --context-switch <c>`  
184     `--state-hashing -I <path-to-CUDA-OM>,`

```

1 #include <...>
2 #define BLOCKS 1
3 #define THREADS 2
4
5 __global__ void kernel(int *A) {
6     A[threadIdx.x + 1] = threadIdx.x;
7 }
8
9 int main() {
10     int *a;
11     int *dev_a;
12     int size = THREADS * sizeof(int);
13     a = (int *) malloc(size);
14     cudaMalloc((void **)&dev_a, size);
15     for (int i = 0; i < THREADS; i++)
16         a[i] = 0;
17     cudaMemcpy(dev_a, a, size, cudaMemcpyHostToDevice);
18     // kernel<<<BLOCKS, THREADS>>>(dev_a);
19     ESBMC_verify_kernel(kernel, BLOCKS, THREADS, dev_a);
20     cudaMemcpy(a, dev_a, size, cudaMemcpyDeviceToHost);
21     for (int i = 0; i < THREADS; i++)
22         assert(a[i] == i);
23     cudaFree(dev_a);
24     free(a);
25     return 0;
26 }

```

Figure 2: Illustrative CUDA code example.

where `<file>.cu` is the CUDA program, `<k>` is the maximum loop unrolling, `<c>` is a context-switch bound, `--state-hashing` reduces redundant interleavings, and `<path-to- CUDA-OM>` is the location of the COM library.

In the mentioned example, ESBMC-GPU detects an array out-of-bounds violation. Indeed, this CUDA-based program retrieves a memory region that has not been previously allocated, *i.e.*, when `threadIdx.x = 1`, the program tries to access `a[2]`. Importantly, the `cudaMalloc()` function’s operational model has a precondition that checks if the memory size to be allocated is greater than zero. In addition, an assertion checks if the result matches to the expected postcondition (line 22). The verification of this program through ESBMC-GPU produces 54 successful and 3 failed interleavings. For instance, one possible failed interleaving is represented by the threads executions  $t_0 : a[1] = 0$ ;  $t_1 : a[2] = 1$ , where  $a[2] = 1$  represents an incorrect access to the array index  $a$ . It is worth noticing that CIVL, ESBMC-GPU, and GKLEE are also able to detect this array out-of-bounds violation, but GPUVerify fails, as it reports a true incorrect result (*i.e.*, a missed bug).

## 5. Experimental Evaluation

In order to evaluate ESBMC-GPU’s precision and performance, a benchmark suite was created, which comprises 20 CUDA kernels from NVIDIA GPU Computing SDK v2.0 [22], 20 CUDA kernels from Microsoft C++ AMP Sample Projects [23], and 114 CUDA-based programs that explore a wide range of CUDA functionalities. In summary, the chosen suite contains 47.4% bug-free and 52.6% buggy benchmarks, which were organized into 5

sets, in order to simplify our discussion, according to the properties it tackles: array bounds (5), assertive statements (7), data-race conditions (17), pointer safety (7), and other specific CUDA functionalities (118). The latter includes `__device__` function calls, general CUDA functions (*e.g.*, `cudaMemcpy`), general libraries in CUDA (*e.g.*, `curand.h`), type modifiers (*e.g.*, `unsigned`), type definitions, and intrinsic CUDA variables (*e.g.*, `uint4`).

The present experiments answer two research questions: (i) How accurate is ESBMC-GPU when verifying the chosen benchmarks? (ii) How does ESBMC-GPU’s performance compare to other existing verifiers? In order to answer both research questions, all benchmarks were verified with 4 GPU verifiers (ESBMC-GPU *v2.0*, GKLEE *v2012*, GPUVerify *v1811*, and CIVL *v1.7.1*), on an otherwise idle Intel Core i7-4790 CPU 3.60 GHz, with 16 GB of RAM, running Ubuntu 14.04 OS. Importantly, all presented execution times are actually CPU times, *i.e.*, only the elapsed time periods spent in the allocated CPUs, which was measured with the `times` system call (POSIX system). An overview of the experimental results is shown in Fig. 3, where *True* represents bug-free benchmarks, *False* represents buggy benchmarks, *Not supported* represents benchmarks that could not be verified, *Correct* represents the percentage of benchmarks correctly verified, and *Incorrect* represents the percentage of benchmarks incorrectly verified (*i.e.*, a verification tool reports an unexpected result).




Figure 3: Experimental evaluation of ESBMC-GPU against other verifiers.

As one may notice, the present experimental results show that ESBMC-GPU reached a successful verification rate of approximately 85%, while GKLEE, GPUVerify, and CIVL reported 72%, 50%, and 35%, respectively. More precisely, ESBMC-GPU correctly detected all data-race conditions present in the benchmarks, which is due to the COM’s multi-threading model that under-approximates the GPU kernels. It also outperformed GKLEE, GPU-



235 Verify, and CIVL, in the verification of array out-of-bounds violations (100%),  
 236 assertive statements (86%), and pointer safety (72%), which is related to  
 237 ESBMC’s capacity to handle arrays and pointers [3]. Furthermore, ESBMC-  
 238 GPU presented the highest coverage rate for specific CUDA functionalities  
 239 (82%) that is once again due to COM, which incorporates specific pre- and  
 240 post-conditions into its verification processes.

241  
 242 **Limitations.** ESBMC-GPU was unable to correctly verify 24 benchmarks,  
 243 which are related to constant memory access (2%), CUDA’s specific libraries  
 244 (*e.g.*, `curand.h`) (4.5%), and the use of pointers to functions, structures, and  
 245 `char` type variables, when passed as kernel call arguments (4.5%). In addi-  
 246 tion, it only reported 3% of incorrect true, which are due to NULL pointer  
 247 accesses, and 1% of incorrect false results, due to partial coverage of the *cu-*  
 248 *daMalloc* function for copies over `float` variables. The remaining verifiers  
 249 (*i.e.*, GKLEE, GPUVerify, and CIVL) were unable to detect mostly data-  
 250 race conditions, assertive statements, and array out-of-bounds violations.  
 251 In addition, they lack support of CUDA specific features, *e.g.*, GPUVer-  
 252 ify does not support the use of the `memset` function nor function point-  
 253 ers and CIVL does not support several CUDA features, such as `atomic`  
 254 functions, `cudaThreadSynchronize`, `threadIdx`, `curand` functions, `dim3`,  
 255 `math_functions`, `uint4`, `__constant__` variables, among others.

256  
 257 **Performance.** MPOR resulted in a performance improvement of approx-  
 258 imately 80%, by decreasing the verification time from 16 to 3 hours, while  
 259 the two-threads analysis further reduced that to 789.6 sec. Although such  
 260 techniques have considerably improved the ESBMC-GPU’s performance, it  
 261 still takes longer than the other evaluated tools: GPUVerify (98.36 sec), GK-  
 262 LEE (105.18 sec), and CIVL (708.52 sec). On the one hand, this is due to  
 263 thread interleavings, which combine symbolic model checking with explicit  
 264 state-space exploration [6]. On the other hand, ESBMC-GPU still presents  
 265 the highest accuracy, with less than 6 seconds per benchmark.

266  
 267 **Availability of Data and Tools.** The performed experiments are based on  
 268 a set of publicly available benchmarks. All benchmarks, tools, and results,  
 269 associated with the current evaluation, are available at [www.esbmc.org/gpu/](http://www.esbmc.org/gpu/).

## 270 6. Conclusions and Future Work

271 ESBMC-GPU marks the first application of an SMT-based context-BMC  
 272 tool that recognizes CUDA directives [6]. Besides, it also applies MPOR,  
 273 two-thread analysis, and state hashing, in order to further simplify verifi-  
 274 cation models and provides fewer incorrect results, compared with GKLEE,

GPUVerify, and CIVL. Indeed, it presents improved ability to detect array out-of-bounds and data race violations.

Future work aims to extend ESBMC-GPU, in order to fully support the verification of CUDA (parallel) streams and events [1]. In addition, more models of libraries will be integrated into COM, with the goal of increasing the coverage of CUDA’s API such as CUDA Driver API, NPP, and cuSOLVER. Finally, we also aim to implement further techniques (*e.g.*, invariant inference via abstract interpretation [24]), in order to prune the state-space exploration, by taking into account GPU symmetry.

## Acknowledgements

This paper is based on research sponsored by the Institute of Development and Technology (INdT) and by the National Council for Scientific and Technological Development (CNPq) under agreement number 475647/2013 – 0.

## References

- [1] J. Cheng, M. Grossman, T. McKercher, *Professional CUDA C Programming*, John Wiley and Sons, Inc., Indianapolis, Indiana, USA, 2014.
- [2] A. Betts, N. Chong, A. Donaldson, S. Qadeer, P. Thomson, *GPUVerify: A Verifier for GPU Kernels*, in: Proceedings of the ACM International Conference on Object Oriented Programming Systems Languages and Applications, ACM, New York, NY, USA, 2012, pp. 113–132, URL <http://doi.acm.org/10.1145/2384616.2384625>.
- [3] L. Cordeiro, B. Fischer, J. Marques-Silva, *SMT-Based Bounded Model Checking for Embedded ANSI-C Software*, IEEE Trans. Software Eng. 38 (2012) 957–974, URL <https://doi.org/10.1109/TSE.2011.59>.
- [4] P. Pereira, H. Albuquerque, H. Marques, I. Silva, C. Carvalho, V. Santos, R. Ferreira, L. Cordeiro, *Verificação de Kernels em Programas CUDA usando Bounded Model Checking*, in: XV Brazilian Symposium on High-Performance Computing (WSCAD-SSC), Florianópolis, Brazil, 2015, pp. 24–35.
- [5] Pereira, P., Albuquerque, H., Marques, H., Silva, I., Carvalho, C., Santos, V., Ferreira, R., and Cordeiro, L. *Verifying CUDA Programs using SMT-Based Context-Bounded Model Checking*. in: Proceedings of the 31<sup>st</sup> Annual ACM Symposium on Applied Computing, ACM, New York, NY, USA, 2016, pp. 1648–1653, URL <http://doi.acm.org/10.1145/2851613.2851830>.

- [6] P. Pereira, H. Albuquerque, I. Silva, H. Marques, F. R. Monteiro, R. Ferreira, L. Cordeiro, *SMT-Based Context-Bounded Model Checking for CUDA Programs*, Concurrency Computat.: Pract. Exper. (2016), URL <http://dx.doi.org/10.1002/cpe.3934>.
- [7] G. Li, G. Gopalakrishnan, *Scalable SMT-based Verification of GPU Kernel Functions*, in: Proceedings of the 18<sup>th</sup> ACM SIGSOFT International Symposium on Foundations of Software Engineering, ACM, New York, NY, USA, 2010, pp. 187–196, URL <http://doi.acm.org/10.1145/1882291.1882320>.
- [8] G. Li, P. Li, G. Sawaya, G. Gopalakrishnan, I. Ghosh, S. Rajan, *GK-LEE: Concolic Verification and Test Generation for GPUs*, in: Proceedings of the 17<sup>th</sup> ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming, ACM, New York, NY, USA, 2012, pp. 215–224, URL <http://doi.acm.org/10.1145/2145816.2145844>.
- [9] M. Zheng, M. Rogers, Z. Luo, M. Dwyer, S. Siegel, *CIVL: Formal Verification of Parallel Programs*, in: Proceedings of the 30<sup>th</sup> IEEE/ACM International Conference on Automated Software Engineering, IEEE Computer Society, Washington, DC, USA, 2015, pp. 830–835, URL <http://dx.doi.org/10.1109/ASE.2015.99>.
- [10] A. Biere, *Bounded Model Checking*, in: A. Biere, M. Heule, H. van Maaren, T. Walsh (Eds.), Handbook of Satisfiability, IOS Press, Amsterdam, The Netherlands, 2009, pp. 457–481.
- [11] C. Barrett, R. Sebastiani, S. Seshia, C. Tinelli, *Satisfiability Modulo Theories*, in: A. Biere, M. Heule, H. van Maaren, T. Walsh (Eds.), Handbook of Satisfiability, IOS Press, Amsterdam, The Netherlands, 2009, pp. 825–885.
- [12] L. Cordeiro, B. Fischer, *Verifying Multi-threaded Software using SMT-based Context-Bounded Model Checking*, in: Proceedings of the 32<sup>nd</sup> ACM/IEEE International Conference on Software Engineering, ACM, New York, NY, USA, 2011, pp. 331–340, URL <http://doi.acm.org/10.1145/1810295.1810396>.
- [13] V. Kahlon, C. Wang, A. Gupta, *Monotonic Partial Order Reduction: An Optimal Symbolic Partial Order Reduction Technique*, in: Proceedings of the 21<sup>st</sup> International Conference on Computer Aided Verification, Springer-Verlag, Berlin, Heidelberg, 2009, pp. 398–413, URL [http://dx.doi.org/10.1007/978-3-642-02658-4\\_31](http://dx.doi.org/10.1007/978-3-642-02658-4_31).

- [14] M. Ramalho, M. Freitas, F. R. Monteiro, H. Marques, L. Cordeiro, B. Fischer, *SMT-Based Bounded Model Checking of C++ Programs*, in: Proceedings of the 20<sup>th</sup> Annual IEEE International Conference and Workshops on the Engineering of Computer Based Systems, IEEE Computer Society, Washington, DC, USA, 2013, pp. 147–156, URL <http://dx.doi.org/10.1109/ECBS.2013.15>.
- [15] F. R. Monteiro, L. Cordeiro, E. B. de Lima Filho, *Bounded Model Checking of C++ Programs Based on the Qt Framework*, in: IEEE 4<sup>th</sup> Global Conference on Consumer Electronics, IEEE Consumer Electronics Society, Washington, DC, USA, 2015, pp. 179–447, URL <https://doi.org/10.1109/GCCE.2015.7398699>.
- [16] M. Garcia, F. R. Monteiro, L. Cordeiro, E. B. de Lima Filho, *ESBMC<sup>QtOM</sup>: A Bounded Model Checking Tool to Verify Qt Applications*, in: 23<sup>rd</sup> International Symposium on Model Checking Software, Springer International Publishing, Cham, 2016, pp. 97–103, URL [http://dx.doi.org/10.1007/978-3-319-32582-8\\_6](http://dx.doi.org/10.1007/978-3-319-32582-8_6).
- [17] F. R. Monteiro, M. Garcia, L. Cordeiro, E. B. de Lima Filho, *Bounded Model Checking of C++ Programs based on the Qt Cross-Platform Framework*, *Softw Test Verif Reliab.* 27 (2017) e1632, URL <http://dx.doi.org/10.1002/stvr.1632>.
- [18] Institute of Electrical and Electronics Engineers Inc., *IEEE Standard for Information Technology-Portable Operating System Interface (POSIX) Base Specifications*, IEEE Std. 1003.1-2008 (Revision of IEEE Std. 1003.1-2004). (2008) c1-3826, URL <https://doi.org/10.1109/IEEESTD.2008.4694976>.
- [19] J. Morse, M. Ramalho, L. Cordeiro, D. Nicole, B. Fischer, *ES-BMC 1.22 (Competition Contribution)*, in: 20<sup>th</sup> International Conference on Tools and Algorithms for the Construction and Analysis of Systems, Springer, Berlin, Heidelberg, 2014, pp. 405–407, URL [http://dx.doi.org/10.1007/978-3-642-54862-8\\_31](http://dx.doi.org/10.1007/978-3-642-54862-8_31).
- [20] National Institute of Standards and Technology, *Secure Hash Standard (SHS)*, Federal Information Processing Standards Publication 180-4. (2015), URL <http://dx.doi.org/10.6028/NIST.FIPS.180-4>.
- [21] J. Morse, *Expressive and Efficient Bounded Model Checking of Concurrent Software*, University of Southampton, PhD Thesis, Southampton, UK, 2015.

- 382 [22] NVIDIA Corporation, *CUDA Toolkit Archive*. <https://developer.nvidia.com/cuda-toolkit-archive>, 2015 (accessed 01.06.2017).
- 383
- 384 [23] D. Moth, *C++ AMP Sample Projects for Download*. Microsoft Corpora-  
385 tion, [blogs.msdn.com/b/nativeconcurrency/archive/2012/01/30/  
386 c-amp-sample-projects-for-download.aspx](https://blogs.msdn.com/b/nativeconcurrency/archive/2012/01/30/c-amp-sample-projects-for-download.aspx), 2012 (accessed  
387 01.06.2017).
- 388 [24] W. Rocha, H. Rocha, H. Ismail, L. Cordeiro, B. Fischer, *DepthK: A k-*  
389 *Induction Verifier Based on Invariant Inference for C Programs*, in: 23<sup>rd</sup>  
390 International Conference on Tools and Algorithms for the Construction  
391 and Analysis of Systems, Springer, Berlin, Heidelberg, 2017, pp. 360–  
392 364, URL [http://dx.doi.org/10.1007/978-3-662-54580-5\\_23](http://dx.doi.org/10.1007/978-3-662-54580-5_23).

393 **Required Metadata**

394 **Current executable software version**

395 Ancillary data table required for sub version of the executable software.

Nr.	(executable) Software metadata description	Please fill in this column
S1	Current software version	2.0
S2	Permanent link to executables of this version	<a href="http://esbmc.org/gpu/">http://esbmc.org/gpu/</a>
S3	Legal Software License	Apache v2.0
S4	Computing Operating System	Ubuntu Linux OS
S5	Installation requirements & dependencies	GNU Libtool; Automake; Flex & Bison; Boost C++ Libraries; Multi-precision arithmetic library developers tools (libgmp3-dev package); SSL development libraries (libssl-dev package); CLang 3.8; LLDB 3.8; GNU C++ compiler (multilib files); libc6 and libc6-dev packages
S6	Link to user manual	<a href="http://esbmc.org/gpu/">http://esbmc.org/gpu/</a>
S7	Support email for questions	<a href="mailto:lucas.cordeiro@cs.ox.ac.uk">lucas.cordeiro@cs.ox.ac.uk</a>

Table 1: Software metadata (optional)

396 **Current code version**

397 Ancillary data table required for subversion of the codebase.

Nr.	Code metadata description	Please fill in this column
C1	Current code version	<i>v2.0</i>
C2	Permanent link to code/repository used for this code version	<a href="https://github.com/ssvlab/esbmc-gpu">https://github.com/ssvlab/esbmc-gpu</a>
C3	Legal Code License	GNU Public License
C4	Code versioning system used	git
C5	Software code languages, tools, and services used	C++
C6	Compilation requirements, operating environments & dependencies	GNU Libtool; Automake; Flex & Bison; Boost C++ Libraries; Multiprecision arithmetic library developers tools (libgmp3-dev package); SSL development libraries (libssl-dev package); CLang 3.8; LLDB 3.8; GNU C++ compiler (multilib files); libc6 and libc6-dev packages
C7	Link to developer documentation	<a href="http://esbmc.org/gpu">http://esbmc.org/gpu</a>
C8	Support email for questions	<a href="mailto:lucas.cordeiro@cs.ox.ac.uk">lucas.cordeiro@cs.ox.ac.uk</a>

Table 2: Code metadata (mandatory)