

# ESBMC-CHERI: Towards Verification of C/C++ Programs for CHERI Platforms with ESBMC

**Franz Brauße**<sup>1</sup>   Kunjian Song<sup>1</sup>   Fedor Shmarov<sup>2</sup>   Rafael Menezes<sup>1</sup>  
Mikhail R. Gadelha<sup>3</sup>   Konstantin Korovin<sup>1</sup>   Giles Reger<sup>1</sup>   Lucas C. Cordeiro<sup>1</sup>

<sup>1</sup>The University of Manchester, UK

<sup>2</sup>University of Newcastle, UK

<sup>3</sup>Igalia, A Coruña, Spain

POCL'24 London, 16/01/2024

# Overview of this talk

- Brief introduction to ESBMC
- Architectural changes necessary to support CHERI-C
- Challenges and solutions
- Conclusions and future work

# Introduction to ESBMC

## Efficient SMT-based Bounded Model Checker

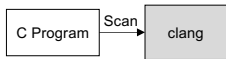
- Open source logic-based automated reasoning for safety and security of C/C++ programs.
  - Apache License 2.0.
  - Used by industries: Intel, Nokia, ARM
  - Core concept: *symbolic execution*.
  - *Concurrency* support.
  - Frontends for *different input languages*.
- **Input:** Program, **Output:** verification *success* or *failure* + witness.
  - *Witness* is a trace of assignments leading to the violation.
- C/C++ frontends translate the Clang AST.

# Architecture of ESBMC

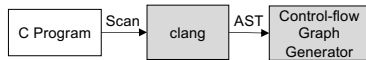
```
graph LR; A[C Program] --> B[ ]; B --> C[Model];
```

C Program

# Architecture of ESBMC



# Architecture of ESBMC



# Architecture of ESBMC

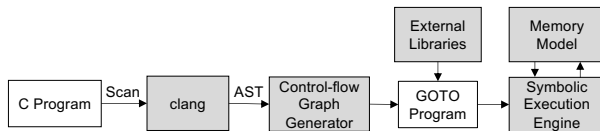


# Architecture of ESBMC

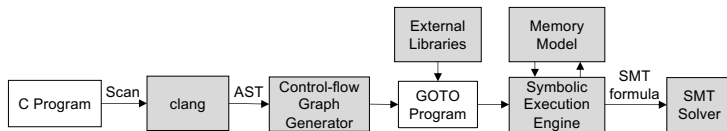




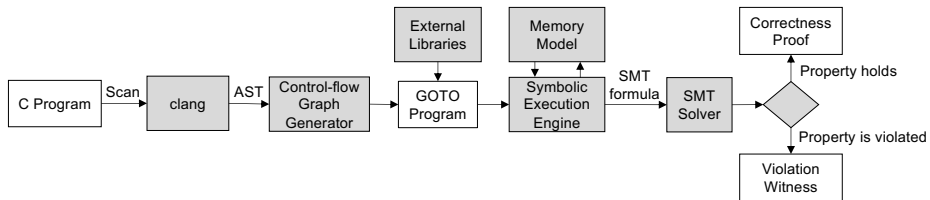
# Architecture of ESBMC



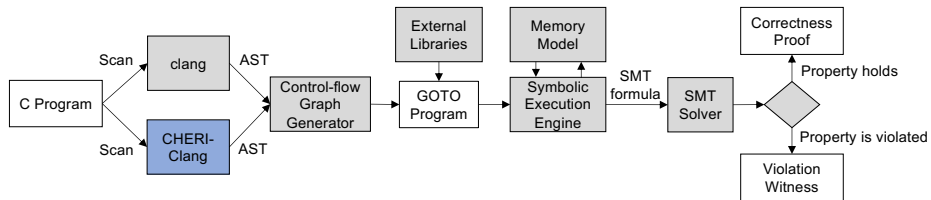
# Architecture of ESBMC



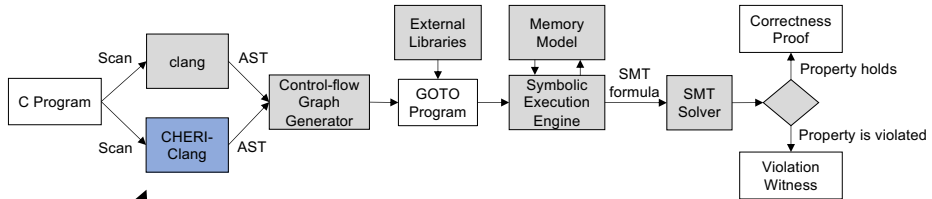
# Architecture of ESBMC



# Architecture of ESBMC

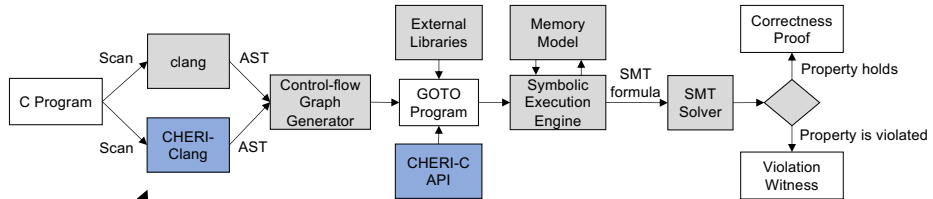


# Architecture of ESBMC



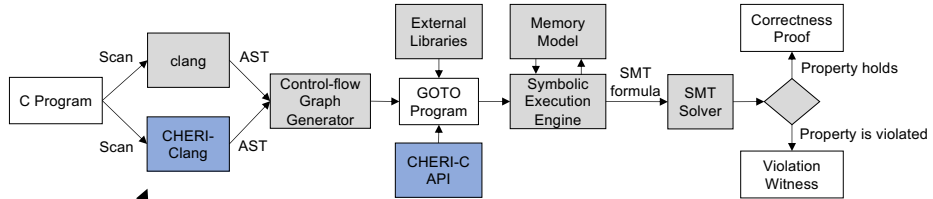
CHERI Clang/LLVM  
compiler

# Architecture of ESBMC



CHERI Clang/LLVM  
compiler

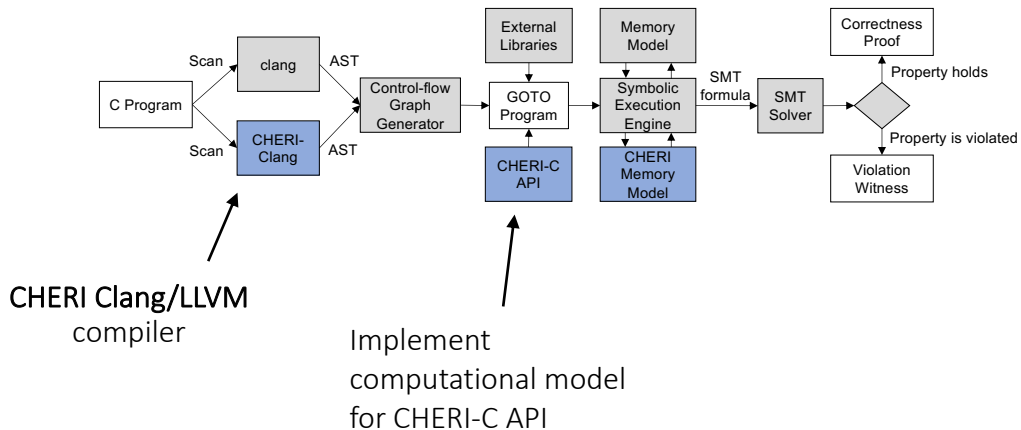
# Architecture of ESBMC



CHERI Clang/LLVM  
compiler

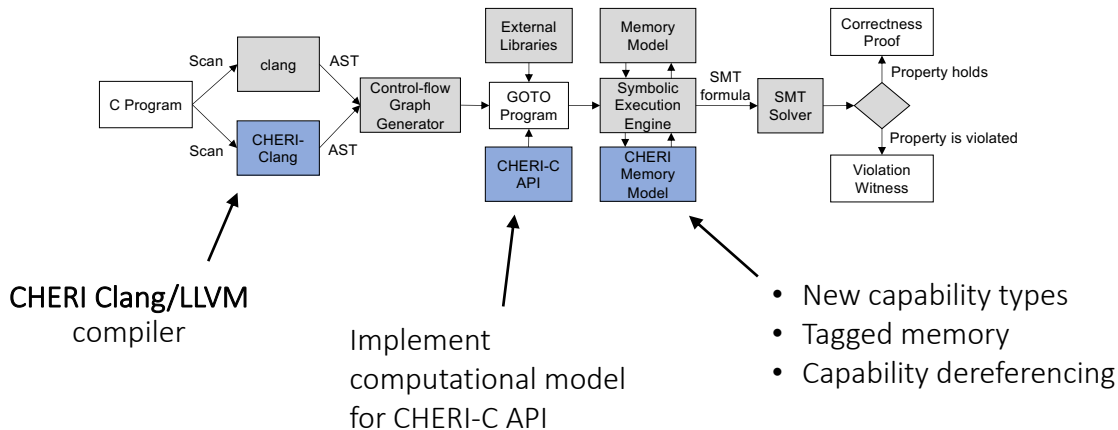
Implement  
computational model  
for CHERI-C API

# Architecture of ESBMC





# Architecture of ESBMC



# ESBMC-CHERI

## Goals:

- both **uncompressed** and **concentrate** capabilities modelled **bit-precise**
- reasoning in **CHERI-BSD** execution environment of capability hardware platforms  
**RISC-V**, **ARM Morello**, (**MIPS**)
- both, **hybrid** and **purecap** mode

# ESBMC-CHERI

## Goals:

- both **uncompressed** and **concentrate** capabilities modelled **bit-precise**
- reasoning in **CHERI-BSD** execution environment of capability hardware platforms  
**RISC-V, ARM Morello, (MIPS)**
- both, **hybrid** and **purecap** mode

Challenges wrt. compression algorithms

# ESBMC-CHERI

## Goals:

- both **uncompressed** and **concentrate** capabilities modelled **bit-precise**
- reasoning in **CHERI-BSD** execution environment of capability hardware platforms **RISC-V**, **ARM Morello**, (**MIPS**)
- both, **hybrid** and **purecap** mode

## Challenges wrt. compression algorithms

## Two directions for verification of CHERI-C programs:

- **Proving** absence of CHERI-exceptions
  - Bit-precise model of in-memory representation of capabilities
  - Base/bounds might differ from those ESBMC already maintains
  - Additional checks for permissions
- **Assuming** absence of CHERI-exceptions
  - Reasoning about unmodelled **external** C functions (libraries) via pointer **provenance**
  - **Simplification** of ESBMC's **spatial memory safety** checks
  - **Optimisation** of **operational models** of `libc` functions

# Challenges for Pointer Provenance

Non-provenance in 'real world'

## 1 International Competition on Software Verification (SV-COMP<sup>1</sup>):

```
void *p = malloc(n);
intptr_t i = p;
free(p);
p = malloc(n);
assert(i != p);
```

Should **fail** verification.

## 2 ARM Realm Management Monitor project<sup>2</sup>:

```
char buffer[n];
intptr_t i = nondet(); /* user-input */
assume(buffer <= i && i < buffer + n);
char *q = buffer + ((char *)i - buffer);
*q;
```

Expected to **work**.

---

<sup>1</sup><https://sv-comp.sosy-lab.org>

<sup>2</sup>Part of Arm Confidential Compute Architecture

# Conclusions

- CHERI-Clang makes frontend changes easy.
- Tagged memory & CHERI-C API using existing functionality.
- Provenance not universally compatible with existing expectations.
- In-memory representation of capabilities and their permissions are difficult.
- On the other hand: assuming CHERI semantics reduces complexity of software verification.